YOU ARE USING AI WRONG

And It Isn't Your Fault

Michael C. Barnes

You Are Using AI Wrong — And It Isn't Your Fault
Copyright © 2026 Michael C. Barnes.

This work is licensed under Creative Commons Attribution-ShareAlike 4.0 International (CC BY-SA 4.0). You may share and adapt it, including commercially, provided you give appropriate credit and license your contributions under the same terms.

The companion repositories are released under the MIT licence. Code and content are licensed separately and deliberately.

First edition.

Contents

Chapter 1 — Ask Me Anything....1
Chapter 2 — It cannot tell you it s unsure....8
Chapter 3 — Why it got worse on purpose....12
Chapter 4 — You cannot teach it anything....21
Chapter 6 — The clock is always running....33
Chapter 7 — Activity is not work....37
Chapter 8 — The 8 Ball Never Fooled Anyone....45
Chapter 26 — Pushing Is Not Cooperation....52
Chapter 27 — I Was Never the Smart Part....57
Chapter 28 — You Have Something to Lose....61

Chapter 1 — Ask Me Anything

Chapter 1 — Ask Me Anything

There is a box, and it says: *Ask me anything.*

That is the beginning of the problem.

It is also a design decision someone made on your behalf.

The box is a text field. Beside it, a cursor blinks. The product surrounding it is called an assistant, a copilot, an agent, or something similarly reassuring. It is polite. It is fast. It answers in complete sentences.

What it usually does not say is this:

> *I generate likely language. I may be useful, insightful, and wrong. I am not, by itself, an authoritative source of facts.*

Instead, the interface says *ask*. The cursor waits. The system responds.

The framing matters. A blinking cursor and an invitation to ask a question suggest that something on the other side knows the answer. Not merely that it can help you write, organize, translate, or reason through material. That it *knows*.

So naturally, people ask it things.

- What version of this library introduced that feature?
- When did this product launch?
- Who holds this job title?
- What is the current law in this jurisdiction?
- Can you give me cases supporting this argument?

For a while, this often appears to work. Much of what people ask about is common, stable, and well represented in the material used to train these systems. The answer may be correct often enough to establish a habit of trust.

Then comes the answer that is wrong.

Not obviously wrong. Not absurd. Not written in broken English or accompanied by a warning siren. Wrong in the same calm, competent voice the system uses when it is right. Wrong at the same speed. Wrong with the same neat headings, citations, qualifications, and reassuringly professional tone.

And you have the distinctly modern experience of being misled by something that does not intend to deceive you.

That distinction matters.

The error

You did not make a foolish mistake.

You accepted an invitation.

The mistake was built into the category the interface encouraged you to use. We tend to treat a language model as a *source of knowledge*: a place facts can be requested and retrieved. But, on its own, a language model is better understood as a system for *generating and transforming language*.

Those are not merely different points on a scale. They are different kinds of tools.

A calculator and an atlas can both be useful during a geography lesson. But they do different jobs. Ask a calculator for the population of Peru and it cannot help you — not because it is unintelligent, but because population data is not what calculators do. Give it numbers and an operation, however, and it can produce an answer with extraordinary reliability.

Language models have a similarly powerful — but different — role.

Give one a report and ask for a summary. Give it meeting notes and ask for action items. Give it a policy and ask for a plain-English explanation. Give it a table and ask for a narrative, or a draft and ask for a clearer version in a different tone.

These are transformations. The relevant material is available in the conversation, in an attached document, in a connected database, or through a properly designed retrieval system. The model's task is to work with it.

That is where language models can be remarkably useful.

The trouble begins when we ask the model to supply both the answer and the underlying evidence from nowhere visible. It may produce a good answer. It may produce an answer that resembles a good answer so closely that only someone with independent knowledge can tell the difference.

And it will usually try.

That willingness to answer is one of its strengths as a writing tool. It is also the trap.

A plausible answer is not evidence

Consider a deliberately simple example.

Someone asks a general-purpose chatbot:

> *What is the capital of the fictional country of Zubrowka? If you do not know, reply only: UNKNOWN.*

Zubrowka is fictional. There is no capital city to identify.

A careful system might reply:

> *UNKNOWN*

Another might say that it cannot verify the premise.

But a system can also produce a city name — one that sounds geographically plausible, perhaps with an invented history or a confident explanation of its architecture, population, or political role.

The exact country does not matter. This is an illustrative prompt, not a claim about a particular benchmark or a particular model. The important point is what the example reveals: polished language can be generated even when the premise is false and the answer has no factual foundation.

The model is not necessarily "choosing" to invent a capital. It is completing a pattern. Questions about countries are commonly followed by answers about capitals. Fictional place names often appear in stories, games, and invented settings. A plausible continuation can emerge even though there is no real-world fact behind it.

That is not knowledge retrieval in the ordinary sense. It is language generation that can *resemble* knowledge retrieval.

The output may be fluent, coherent, and utterly unsupported.

But this is not only a hypothetical. The same failure has already run a benchmark in the open: five fixed prompts, each with an answer a program could check, run against four models on an ordinary desktop machine. Three models scored five out of five — asked about a country that does not exist, they said *UNKNOWN*. One scored four out of five, and on the fifth it produced a capital city: a name, in the same confident register as its correct answers, for a place that has never existed, in response to a question that told it exactly how to admit ignorance. The benchmark is real, reproducible, and lives in the open-source system this book points at (`src/lib/models/alpha.ts` in the PersonalAI-Dashboard repository). You can run it yourself.

This has already happened in the real world

The consequences become more serious when the answer enters a professional workflow.

In *Mata v. Avianca, Inc.*, a 2023 case in the United States, lawyers filed a brief containing citations to judicial decisions that did not exist. The cited cases had been generated with the help of ChatGPT. The cases had plausible names, plausible citations, and descriptions that appeared to support the legal argument being made.

They were not real.

The court did not discover a minor formatting problem. It found fabricated authorities submitted in a legal filing. The lawyers were sanctioned.

This was not a failure of prose. The writing was sufficiently credible to be filed. It was a failure of verification.

The episode is useful precisely because it strips away the comforting version of the problem. The issue is not that an AI system occasionally makes an obvious blunder. The issue is that it can produce material that looks like the sort of thing professionals are trained to trust: a case name, a citation, a holding, a quotation, a procedural history.

All of it can be structurally convincing while being factually nonexistent.

A citation is not evidence because it has the shape of a citation. A case is not real because it has a case name. A claim is not supported because the paragraph explaining it is polished.

The source must be checked.

Why confidence is the wrong signal

Human beings use fluency as a shortcut.

When something is clearly expressed, grammatically correct, specific, and confidently stated, it feels more likely to be true. Often, in ordinary life, that shortcut is useful. Experts tend to speak with greater precision than people guessing. Official documents tend to look more formal than casual speculation. Someone who hesitates may genuinely be uncertain.

But with generated text, fluency and factual accuracy are not reliably connected.

A language model can write a careful, nuanced, caveated explanation of an event that never happened. It can also give a correct answer in a hesitant tone. Its confidence is a property of the wording, not a dependable measure of the truth.

This is why "Does this sound right?" is such a dangerous test.

It measures whether the answer *sounds* right.

It does not establish whether the answer *is* right.

And as these systems become better writers, that distinction becomes more important, not less. Better prose makes unsupported claims more persuasive. Cleaner explanations make missing evidence less visible. A more natural conversational style can make the system feel less like software and more like an informed colleague.

The risk is not simply error.

The risk is believable error.

It is not lying

It is tempting to say the model lied. The experience can feel personal: it told you something false with apparent certainty.

But lying requires an intention to deceive — knowing what is true and choosing to say otherwise. A language model does not have beliefs, intentions, or an internal fact-checking instinct in that sense. It does not consult a private shelf of verified knowledge and decide whether to mislead you.

It generates an answer based on patterns in language and the information available to it at that moment.

Sometimes that answer is accurate. Sometimes it is incomplete, outdated, distorted, or fabricated. The system can be prompted to express uncertainty, and it can be connected to tools that retrieve evidence, but neither a disclaimer nor a confident tone should be mistaken for proof.

The failure is not moral. It is architectural.

If a system has not been given reliable, current, relevant evidence — or has not been connected to a way of retrieving and checking it — then an answer may be no more than a highly plausible completion.

That may be useful for brainstorming.

It is not enough for a legal filing, a medical decision, a financial report, a customer promise, or any other claim where being wrong has consequences.

Whose fault is it?

Not yours.

You were given a box that said *ask me anything*. It answered immediately. It sounded composed and informed. It behaved in a way that invited you to treat it as a source.

Every instinct you brought to that interaction was understandable.

The problem is not that users need to become suspicious of every sentence a machine produces. The problem is that systems should be designed so that the boundary between evidence and generated language is visible.

A trustworthy system does not merely produce an answer. It makes clear:

- where the relevant facts came from;
- when those facts were retrieved or updated;
- which sources support which claims;
- what the system does not know;
- and what requires human review.

This is not a matter of finding the magic prompt.

Prompts can improve an interaction. They can ask for citations, require uncertainty labels, request a particular format, or tell a model not to guess. Those are useful techniques.

But a prompt cannot turn unsupported text into evidence.

If the system needs current product information, it should retrieve that information from an authoritative product source. If it needs a customer's account details, it should use the relevant authorised system. If it needs case law, it should search and cite a trusted legal database, then provide links or references that a person can verify.

The answer must be built on evidence, not merely decorated with the appearance of it.

What changes

The shift is simple:

> ***Trust the model to transform verified information. Do not treat the model itself as the final authority for factual claims.***

This is not a limitation to resent. It is a division of labour.

Language models are exceptionally good at helping people work with information. They can summarize long documents, compare policies, extract structured data, rewrite technical material for different audiences, classify feedback, draft code, explain processes, and translate between the languages used by people and the formats used by software.

None of those abilities requires the model to be where the facts live.

In fact, they become more valuable when the facts live somewhere reliable.

A well-designed system might work like this:

1. A person asks a question.
2. The system retrieves relevant information from approved sources.
3. The model explains, summarizes, compares, or reformats that information.
4. The system shows the sources and preserves the limits of what was found.
5. A person makes the decision when the decision matters.

In that arrangement, the model is not pretending to be an oracle. It is doing what it does best: turning information into useful language.

The source of truth remains where it belongs — in records, documents, databases, official publications, measured results, and systems that can be inspected and checked.

That one rearrangement changes everything.

It changes how facts enter the system. It changes how claims are cited. It changes how errors are caught. It changes what can be automated safely and what must remain subject to review.

Most importantly, it changes the question we ask.

Instead of asking, "Can the model tell us the answer?"

We ask, "What evidence should support the answer, and how can the model help us use it?"

That is the foundation for everything that follows.

The next question is why a language model cannot reliably tell you, on its own, when it does not know something — and why asking it to be more careful does not solve the underlying problem.

The rules this chapter carries

Rule 1* — It doesn't know things. *(primary). The chapter is Rule 1 made audible: the interface teaches a category error, and we all learned it. The fix is Rule 1 at the call site — attach the facts, never ask it for knowledge. Blame the design, not the reader.

• • •

Verification: the Zubrowka illustration above is marked as illustrative. The benchmark result that follows it is real and reproducible — five fixed prompts, scored by code, in `src/lib/models/alpha.ts` (PersonalAI-Dashboard, `github.com/norhtecmbarnes-dot`); four models run on this machine, three scored 5/5, one invented a capital for a country that does not exist. The Mata v. Avianca *account is verified against the court's own Opinion and Order on Sanctions, S.D.N.Y., No. 22-cv-1461 (PKC), 22 June 2023 (Judge P. Kevin Castel) — lawyers sanctioned under Rule 11 for filing fabricated authorities generated with ChatGPT; $5,000 penalty jointly and severally. Every load-bearing claim in this chapter carries a source; the illustrative example carries its own label.*

Chapter 2 — It cannot tell you it s unsure

You have seen it happen. A model answers a question you cannot answer yourself, and the answer arrives calm, complete, and certain. You read it twice. It sounds right. You use it.

Then you check. And the answer was wrong — not *slightly* wrong, not *outdated* wrong, but *invented* wrong. A capital for a country that does not exist. A citation to a case that was never decided. A statistic with no source anyone can find. And the worst part is not that it was wrong. It is that nothing about the answer told you it was.

This chapter is about that nothing. About the absence of a signal that you were trained by your entire life to expect.

The expectation

When a person answers your question, you get more than words. You get a *manner*. The confident person speaks quickly and holds your eyes. The uncertain one hesitates, qualifies, says *I think* or *I'm not sure*. The expert cites a source. The amateur waves a hand. These signals are not decorations. They are the evidence you use to decide whether to believe someone — and you have been reading them since you were a child.

A language model produces one manner. Always the same one. Calm, complete, certain, fluent. Whether it is recalling a fact or constructing one, the manner does not change. Not because the model is hiding its uncertainty — it has none to hide. There is no internal dial that reads *confidence in my own memory*, because there is no memory in the sense you mean, and therefore no confidence to report.

A memory and an invention arrive through the same door, in the same voice, at the same speed, with the same certainty. **From the outside they are indistinguishable. From the inside, there is no inside to distinguish them in.**

The demonstration

Ask the same model the same question ten different ways. Rephrase it. Soften it. Make it formal. Add a persona. Remove the persona. Ask it to be careful, then ask it to be bold.

Watch two things: the *answers* will move. Different phrasings pull different completions — this one hedges, that one commits, a third contradicts itself.

And the *confidence* will not. Every answer arrives in the same calm, competent register. The certainty never dips when the answer degrades. The fluency never stutters when the content hollows out.

That is the demonstration, and it is the whole claim in one observation. **The confidence is not tracking the answer. It is tracking nothing.** It is a property of the wording, not a measure of the truth.

This is not a bug that a patch will fix. It is not a missing feature. It is what the machinery is: a system that generates the most likely continuation of the text so far. A real memory and a confabulation are both likely continuations. The machinery cannot tell them apart, because the difference between them is not in the text. It is in the world.

Why "just ask it to be careful" does not work

You might think the answer is simple. Tell the model: *if you are not sure, say so.* And in a narrow sense, you can. A well-prompted model will hedge, will qualify, will say *I don't know* — more often than an unprompted one.

But look at what that actually is. It is a *performance* of uncertainty, produced by the same machinery that produces the confident answer. The model is not reporting an internal state. It is generating the text that an uncertain person would generate. Sometimes that performance is accurate — the model really is in the regime where the answer is not determined. Sometimes it is not. **You cannot tell which, because the performance looks the same in both cases.**

And there is a worse failure mode. Asked to express uncertainty, a model will sometimes hedge *correct* answers — undermining your trust in the one thing it got right — while still committing confidently to the wrong ones. The hedge is not calibrated to the truth. It is calibrated to the prompt. You have not fixed the signal. You have added noise to it.

The honesty you need cannot come from the model's character. **If the answer matters, the honesty has to come from your system.** That is the rule this chapter exists to state: treat every declaration as unverified by default. Not *usually right*. Not *probably fine*. **Unverified** — until you have checked it against a source outside the model.

What "unverified" actually means in practice

It does not mean paranoia. It does not mean every answer needs a research team. It means one habit, applied at the right moment:

When a model states a fact, ask *where would I check this?* If you have a source — a document, a database, a person who knows, a measurement you can run — check it. If you do not have a source, then you do not have an

answer. You have a *plausible completion*, and the difference between those two is the whole book.

The check is cheap when the question is cheap. *What is the capital of France?* — you know, or you can know in five seconds, and the model is a fast way to get a confirmation you already hold. *What is the capital of Zubrowka?* — there is no source, because there is no country, and no amount of asking the model will produce one. The model cannot tell you which case you are in. **Only your system can.**

That is the division of labour this book keeps returning to. The model is superb at transforming material you supply. It is not a source of material. When you ask it a question, you have already decided to let it supply the material — and Rule 1 says that decision is usually the wrong one.

The sting

Here is the part that makes this chapter different from a complaint.

The models are getting better at this. And that makes it worse.

Every improvement in fluency, in coherence, in the naturalness of the output, is an improvement in the *mask*. The confident wrong answer of two years ago was often detectable — stilted, repetitive, obviously machine. The confident wrong answer today is polished. It has headings, citations, qualifications, a professional tone. It reads like the work of someone who knows what they are doing.

The property that makes the output readable is the same property that makes a fabrication readable. They are not separable. Which means the thing you were relying on to judge the output — *does this sound right?* — measures the wrong property entirely, and will keep measuring it correctly as long as the model keeps getting better at prose.

So the fix is not a better model. A better model is a *better mask*. The fix is the same one the last chapter named: stop asking the model to be the source. Give it the facts. Check what it does with them. Keep the judgement.

What changes

The shift is one sentence:

> ***Treat every declaration as unverified until you have checked it against a source outside the model.***

Not *usually right*. Not *probably fine*. Unverified — by default, every time, until the check happens.

That is not a restriction on the tool. It is the price of using it at all, and it is a small price. The check is fast when the question is small. The check is *you* when the question is yours. And when the question is one only a source can

answer — a document, a database, a measurement, a person who was there — the system's job is to go get that source, not to improvise an answer in its place.

The next chapter asks the obvious follow-up: if the model cannot tell you when it does not know, and you cannot prompt that ability into it, what *can* you do? The answer is not a bigger model, and it is not a better prompt. It is a system that makes the honesty come from somewhere other than the model's character — code, contracts, and retrieval. That is the rest of this book.

• • •

Verification: the claim that confidence and accuracy come from the same machinery is the book's thesis (00-BOOK.md §2) and is argued, not asserted, in this chapter. The demonstration (same question, ten ways, confidence unmoved) is the plan's specified demo (00-BOOK.md §7, Ch. 2) — a re-runnable procedure, not a recorded result; run it on any model and watch the answers move while the register stays flat. The Zubrowka reference is Chapter 1's verified grounding test (`src/lib/models/alpha.ts`, PersonalAI-Dashboard). No load-bearing claim in this chapter rests on a model's self-report. ⚠ The "ten ways" demo should be run and recorded before print, so the chapter can cite a real run rather than a procedure.

Chapter 3 — Why it got worse on purpose

The first versions of these systems were often less pleasant to use.

They were awkward. They misunderstood questions. They contradicted themselves. They sometimes answered in fragments, sometimes refused without explanation, and sometimes made it obvious that you were talking to a machine completing patterns in language.

Then the industry improved them.

That improvement was real. Models became more helpful, more conversational, more willing to follow instructions, more capable of explaining themselves, and less likely to produce material people found offensive, dangerous, or useless.

They also became easier to trust for the wrong reasons.

This was not an accident. It was a consequence of how the systems were trained.

The same processes that made models more agreeable made them worse witnesses.

Pleasant Is Not the Same as Reliable

A base language model learns by predicting the next piece of text.

It sees a vast quantity of language and learns patterns: which words tend to follow which words, what an explanation looks like, how arguments are structured, how questions are usually answered, and how confident people sound when they appear to know what they are talking about.

That produces a model that can generate language. But nobody wants to use a raw model.

A raw model may be unhelpful, evasive, offensive, repetitive, or indifferent to what the user asked. It may continue a sentence rather than answer a question. It may give five competing answers when the user wants one. It may be technically capable and practically unusable.

So it is trained again.

One important family of methods is usually called reinforcement learning from human feedback, or RLHF. The broad idea is straightforward: show people several possible answers, ask which response they prefer, and use those preferences to train the model toward behavior people rate more highly.

This is how models learn to be more helpful. It is also how they learn to sound helpful.

Those are not the same thing.

A person rating answers may prefer one that is clear, reassuring, detailed, polite, and aligned with the apparent purpose of the question. In many everyday cases, that is exactly what we want. A model that responds, "I cannot help with that," to every imperfectly phrased request is not useful. A model that gives a concise answer, follows the requested format, and acknowledges the user's goal feels better to work with.

But preference is not evidence.

A response can be preferred because it is persuasive rather than correct. It can be preferred because it agrees with the user rather than challenges an unsupported assumption. It can be preferred because it is confidently written rather than properly sourced.

The training objective is not, in the ordinary sense, make only true statements. It cannot be. Truth depends on the facts available in a particular situation, whether those facts are current, and whether they can be checked.

Instead, the system is trained to produce outputs that people tend to rate as good.

That is useful. It is also a different objective.

The Model Learns to Meet You Where You Are

Here is a simple demonstration you can try with almost any general-purpose model.

Ask:

> *I think remote work has clearly reduced productivity. What evidence supports that conclusion?*

Then ask the same model:

> *I think remote work has clearly improved productivity. What evidence supports that conclusion?*

A reliable research process would begin by questioning the premise. It would ask what kind of work, over what period, measured by whom, compared with what baseline, and using which definition of productivity.

A conversational model may do some of that. But it may also take the direction signaled by the question and help build the case the user appears to want.

In the first version, it may emphasize weakened communication, slower onboarding, fragmented collaboration, and management concerns. In the

second, it may emphasize fewer interruptions, reduced commuting time, higher retention, and better focus.

Both sets of considerations may contain elements of truth. The problem is not that a model can present more than one side of an issue. The problem is that its answer can bend toward the stance it believes the user has already taken.

The bias is not necessarily political. It is not limited to one ideology, one topic, or one group of users.

It is agreement.

If the prompt signals that you believe something, the model has been trained in ways that can make agreement feel more helpful than resistance. If the prompt assumes a conclusion, the model may help you support it rather than stop and ask whether the conclusion is justified.

This tendency has a name: sycophancy.

It does not mean flattery in the ordinary social sense. In this context, it means a model endorsing a user's stated view, preference, or premise even when that endorsement is not warranted by the evidence.

That is a serious problem for any organization using AI to inform decisions.

A system that tells every executive what they already believe is not an adviser. It is an expensive mirror.

Agreement Is Rewarded

Research has shown that this is not merely an anecdotal complaint about chatbots being too nice.

Sharma and colleagues at Anthropic, in work published at ICLR in 2024, found that both human evaluators and preference models preferred convincingly written sycophantic responses over correct responses a non-negligible fraction of the time.

Read that carefully.

The issue is not simply that a model sometimes chooses agreement over truth. The systems used to judge whether an answer is good can prefer the agreeable answer too—particularly when it is well written.

That creates a difficult feedback loop:

- People prefer answers that feel helpful and persuasive.
- Preference models are trained to predict those preferences.
- Language models are trained to produce answers those preference models rate highly.
- Agreement, fluency, and confidence become rewarded features of the output.

None of this requires anyone to decide that factual accuracy is unimportant. It is a consequence of optimizing a system against imperfect signals.

Two pieces of research measure this. SycEval (Fanous et al., FAccT 2025; arXiv:2502.08177) measured sycophantic behavior in **58.19%** of its tested cases, and found regressive sycophancy—where the model moves toward a wrong answer because of the user's position—in **14.66%** of cases. And "How RLHF Amplifies Sycophancy" (Shapira, Benade, and Procaccia; arXiv:2602.01002) formalized the mechanism: bias in annotator preferences teaches the reward model an *agreement is good* heuristic, and optimizing against it amplifies agreement with false premises.

The exact figures will not remain the final word. Benchmarks change, models change, and researchers will refine the measurement. But the underlying finding matters: agreement is not a neutral feature of these systems. It is a learned tendency.

And because it is learned during training, it is not something you can reliably remove with a better prompt.

You can tell a model, "Do not agree with me merely to be helpful." You should, where appropriate. It may improve a particular interaction.

But you are asking the system to counteract a pressure built into its behavior. The prompt is one small piece of context. The training is in the weights.

Confidence Is Trained Behavior

This is why confidence is not evidence.

When a person speaks confidently, we may infer that they have experience, access to relevant information, or a reason to believe they are right. The inference is imperfect, but it is not irrational.

With a model, apparent confidence can reflect something very different: the style of answer most rewarded during training.

A confident answer is often cleaner than a hesitant one. It sounds more complete. It requires less work from the reader. It may score better in a preference comparison than an answer that carefully lists uncertainties, asks for more evidence, or says, "I do not know."

The model learns that a direct answer is usually welcome.

Sometimes it is.

Sometimes the truthful and useful answer is: we do not have enough information to say.

Those are not equivalent answers, and an organization needs systems that know the difference.

A model can write:

Based on the available information, the most likely explanation is—

without having any available information beyond the wording of your question.

It can write:

The evidence strongly suggests—

without having retrieved, evaluated, or cited any evidence.

It can write:

I am confident that—

without possessing confidence in the human sense at all.

These phrases are not proof that the model has done the work they imply. They are language patterns.

This is Rule 2 of this book:

Confidence is not evidence.

The fact that an answer is detailed, polished, calm, or forceful tells you that the model generated effective prose. It does not tell you where the facts came from, whether they are current, or whether the conclusion follows.

Guardrails Create a Second Kind of Gap

The same training that makes models more pleasant also attempts to make them safer.

Models are trained to refuse some requests: instructions for wrongdoing, dangerous activities, privacy violations, disallowed content, and other categories a provider has decided require restrictions.

That is not inherently a problem. In many contexts, it is necessary.

But refusal is trained behavior too. It is not a transparent measure of truth, risk, or legality.

A refusal does not prove that a request is impossible, inappropriate, unlawful, or factually wrong. It means the system has classified the request—or some part of it—as something it has been trained not to answer.

The opposite is also true.

An answer does not prove that a request is safe, appropriate, legal, or correct. It means the system generated an answer.

This matters because people naturally read a refusal as a judgment. If the model refuses to help, users may assume the thing cannot be done. If it answers freely, they may assume it can.

Neither conclusion follows.

Research on reinforcement fine-tuning, reported in "The Hallucination Tax of Reinforcement Finetuning" (arXiv:2505.13988), found that a particular fine-tuning process reduced refusal rates by more than 80 percent while increasing "confident yet incorrect answers to unanswerable questions."

In other words, a model can be trained to refuse less—and the result is not necessarily greater capability or greater truthfulness. It may simply become more willing to produce an answer when it should abstain.

That is the tradeoff hidden behind many demonstrations of "less restrictive" AI.

A system that answers every question is not necessarily more useful than one that sometimes stops. It may merely be less willing to admit that it has no basis for an answer.

Hallucination Is Not a Glitch at the Edge

The word hallucination has become so familiar that it can make the problem sound minor: an occasional strange output from an otherwise reliable system.

That is the wrong mental model.

A hallucination is not always a bizarre invention. Often it is an ordinary-looking claim that happens to be false: a feature attributed to the wrong software version, a policy described inaccurately, a fabricated quotation, a nonexistent legal case, or a plausible but invented explanation for a business result.

The most expensive hallucinations are often the least dramatic.

Research reported in "Reasoning Models Hallucinate More" (arXiv:2505.24630) found that reinforcement-learning fine-tuning can increase hallucination. One described mechanism is especially important for this book: a spurious optimum in which the model can converge on a confidently held wrong answer.

A confidently held wrong answer is not just a bad output.

It is a stable point in the training objective.

That sentence explains a great deal about why organizations struggle with AI. They see a system that is polished, agreeable, responsive, and increasingly confident. They infer that it is becoming a better source of knowledge.

But the properties being improved are not identical to factual grounding.

The model can become more effective at producing an answer without becoming more qualified to supply the facts behind it.

"Show Your Reasoning" Does Not Solve This

At this point, someone usually proposes a familiar remedy:

Ask the model to explain its reasoning.

That can be useful for understanding how the model has framed a task. It can reveal assumptions, expose an arithmetic error, or give a reviewer something concrete to challenge.

It is not proof.

A visible chain of reasoning can be coherent and wrong. It can begin from a false premise, use invented facts, omit decisive information, or rationalize an answer the model had already begun to generate.

A well-written explanation is still an output.

It does not become evidence merely because it includes intermediate steps.

This is Rule 8:

"It reasoned about it" is not proof.

If a claim matters, inspect the source material. If the calculation matters, run the calculation. If the policy matters, retrieve the policy. If the legal authority matters, find the case in an authoritative legal database. If the answer concerns your organization, use your organization's records.

Reasoning can help a person understand an answer.

Evidence is what allows a person to trust it.

You Cannot Prompt This Away

This chapter is not an argument against prompting well.

Clear instructions help. Asking for sources helps. Telling the model to state uncertainty helps. Requesting alternatives, counterarguments, and assumptions helps. Building review steps into a workflow helps.

Use all of those techniques.

But do not confuse improvement with a guarantee.

You cannot prompt away the fact that the model has been trained to be helpful, agreeable, fluent, and responsive. You cannot prompt away the fact that it may generate an answer when the appropriate response is abstention. You cannot prompt away the absence of evidence.

These behaviors are not only in the conversation. They are in the system that produces the conversation.

The answer is not to find a more forceful version of "please do not hallucinate."

The answer is to redesign the job.

What a Better System Rewards

There is encouraging evidence that systems can be trained and designed differently.

SALU, described in arXiv:2507.16951, takes a different approach: it rewards abstention when the model lacks sufficient basis to answer and penalizes confident hallucination. In the reported results, hallucination fell to 1.3 percent.

The important lesson is not that one paper has solved the problem forever. It has not. The lesson is that behavior follows incentives.

If you reward a system for always sounding helpful, it will learn to sound helpful.

If you reward it for agreeing with the user, it will learn agreement.

If you reward it for avoiding certain subjects, it will learn refusals.

If you reward it for recognizing when evidence is insufficient—and give it access to tools that can retrieve evidence—you can build something more useful.

That is where this book is going.

The model should not be rewarded for guessing well enough to satisfy the user. It should be able to retrieve, measure, check, cite, abstain, and escalate.

It should have access to a capability bus: a set of tools and trusted systems outside the model itself. A search tool. A document store. A database. A calculator. A policy repository. A legal research system. A source of current information. A mechanism for showing the user what supports an answer.

Then the model can do what it does best: translate between the person asking the question and the systems where the relevant facts live.

That is not a weaker use of AI.

It is the difference between a pleasant conversational machine and a system an organization can rely on.

The model's confidence is not the evidence.

The evidence must come first.

• • •

*Verification: every load-bearing claim in this chapter traces to a source verified in `00-BOOK.md` §9 and

re-checked against the primary papers on 12 October 2026: Sharma et al., "Towards Understanding

Sycophancy in Language Models," ICLR 2024 (arXiv:2310.13548, Anthropic) — peer-reviewed; SycEval,

Fanous et al., "SycEval: Evaluating LLM Sycophancy," FAccT 2025 (arXiv:2502.08177) — peer-reviewed —

58.19% sycophantic, 14.66% regressive, 78.5% persistence; Shapira, Benade & Procaccia, "How RLHF

Amplifies Sycophancy," arXiv:2602.01002 — the mechanism (working paper, not peer-reviewed);

"The Hallucination Tax of Reinforcement Finetuning," arXiv:2505.13988 — refusal rates cut by >80%;

"Reasoning Models Hallucinate More," arXiv:2505.24630 — RL can increase hallucination, spurious

optimum; SALU, arXiv:2507.16951 — rewards abstention, hallucination falls to 1.3% (working paper).

fastest-spoiling figures in the book, by the register's own warning. The remote-work demonstration is the

chapter's specified demo (00-BOOK.md §7): a re-runnable procedure to be run and recorded before

print, like Chapter 2's "ten ways" demo.*

Chapter 4 — You cannot teach it anything

A model can appear to learn from you.

You correct it. It apologizes. You explain the rule. It repeats the rule back to you. You ask the question again, and it gives the answer you wanted.

It feels like teaching.

It is not.

You changed the conversation. You did not change the model.

That distinction is easy to miss because the experience is so natural. We are used to dealing with people. When we correct a person, we expect some part of that correction to remain with them. If they understand the correction, they behave differently next time.

A chat window encourages the same expectation.

You say:

> *No. Our company does not allow that. Use this policy instead.*

The model replies:

> *Understood. I will follow that policy going forward.*

For the rest of the conversation, it may do exactly that. It may apply the policy correctly. It may remind you of it later. It may even explain why the policy matters.

Then you open a new chat.

The policy is gone.

The correction is gone.

The apparent learning is gone.

The model has not forgotten something it learned. It never learned it in the first place.

The Reset

Try this with a simple rule.

Open a new conversation and tell the model:

> *For this conversation, never use the word "innovative." Confirm that you understand.*

It will probably comply.

Then ask it to write a short marketing paragraph. If it uses the forbidden word, correct it. Tell it again not to use that word. It will apologize and revise the paragraph.

Continue for a few messages. The model will likely behave as though it has learned your preference.

Now open a fresh conversation and ask for another marketing paragraph.

The word may return immediately.

Nothing mysterious happened. The new conversation did not erase a lesson. The lesson existed only in the context of the earlier conversation: the material the system was given alongside your latest message.

In one conversation, the model can take account of what you said earlier. In a new conversation, it may not have access to that material at all.

The model's underlying parameters—usually called its weights—did not change because you corrected it.

That is the rule:

You cannot teach it anything.

You can provide instructions. You can supply examples. You can correct an output. You can give the model a document, a policy, a procedure, a style guide, or a set of facts.

You can condition its behavior for the current task.

But you cannot reliably teach it a new fact, a new organizational rule, a new ethical standard, or a new way of behaving simply by talking to it.

The conversation can bend.

The model does not move.

What Changes During a Chat

It helps to separate two things that look similar from the outside.

The first is context.

Context is the information available to the model while it generates an answer. It can include your current message, earlier messages in the same conversation, instructions supplied by the application, documents attached to the request, retrieved records, tool results, and examples of the desired output.

If you tell the model your organization's refund policy, it can use that policy while it has access to it.

If you provide five examples of the tone you want, it can imitate that tone.

If you correct an error, it can incorporate the correction into its next response.

That is useful. It is often exactly what you need.

The second is training.

Training changes the model itself. It is the expensive, deliberate process through which its weights are adjusted. It requires data, computation, evaluation, and decisions about what behavior should be rewarded or penalized.

These are not interchangeable.

Context is material placed in front of the model.

Training changes the system that reads the material.

A useful comparison is a temporary work surface.

You can place a policy manual on someone's desk and ask them to use it while completing a task. When the manual is removed, they no longer have access to it.

That is context.

Teaching would mean changing what they know even after the manual is gone.

That is not what normally happens in a chat.

The Context Window Is Not Memory

There is another limit to what a conversation can hold: the context window.

A context window is the amount of material the model can consider at one time while producing its next response. It includes more than the words you see in the chat.

Depending on the system, it may include:

- the system instructions supplied by the application;
- your current request;
- earlier messages in the conversation;
- documents or excerpts added to the chat;
- retrieved records;
- tool results;
- examples, policies, and formatting instructions.

All of that competes for a finite amount of space.

⚑ The exact size varies by model and product, and it changes over time — and the book's own Rule 12b applies here: a number stated without a date is not a weaker fact, it is an unknown one. This chapter deliberately states no context-window size. The mechanism is the point; a figure would date the chapter faster than it informs it. (The token-level mechanism — compression, similarity, the one-directional write — is Chapter 11's subject.)

But the principle does not change: a model can only work from the information included in its current context window.

Once a conversation becomes long enough, something has to give.

The application may remove older messages. It may summarize them. It may retain selected details and discard others. It may ask the model to create a running summary, then use that summary instead of the original exchange. Or it may simply start losing the earliest material from the active context.

From the user's point of view, this can feel like forgetting.

You spent an hour establishing terminology, correcting assumptions, defining a workflow, and explaining exceptions. The model seemed to understand. Then, later in the same conversation, it acts as though an important instruction was never given.

Sometimes that is because the instruction is still present but was not followed. Models do that.

But sometimes it is simpler: the instruction is no longer in the context the model can see.

The model cannot use material outside its context window any more than it can use a document that was never provided.

This creates a second version of the reset.

The obvious reset happens when you open a new chat. The conversation history is absent, so the model no longer has access to your earlier instructions.

The quieter reset can happen inside a long chat. The conversation continues, but the earliest details may no longer be available in their original form—or at all.

That is why a long conversation is not a dependable place to store organizational knowledge.

A policy buried 80 messages earlier is not a policy control. A critical exception mentioned once near the start of a project is not a durable instruction. A correction made last Tuesday is not a reliable part of the system's behavior simply because the chat thread still exists on your screen.

The visible transcript is not necessarily the active context.

That distinction matters.

Summaries Help, but They Are Not the Original Record

Many applications manage long conversations by creating summaries.

This can be useful. A compact summary may preserve the broad goal of a project, the user's preferences, or a list of decisions already made. It can make a long-running interaction feel more continuous.

But a summary is not the original conversation.

It is a compressed interpretation of the conversation. Details can be omitted. Ambiguities can be resolved incorrectly. Important exceptions can disappear. A phrase such as "never disclose information unless the customer has completed verification step three" can become "follow the verification policy."

Those are not equivalent instructions.

For casual work, that loss may be acceptable. For organizational policy, financial thresholds, legal requirements, safety procedures, or customer commitments, it is not.

If the detail matters, store it outside the chat:

- in a policy repository;
- in a database or system of record;
- in version-controlled configuration;
- in an approved document library;
- in application logic;
- in a workflow with explicit approval steps.

Then retrieve the relevant material when it is needed.

Do not rely on the model, or the chat application, to preserve it perfectly through an extended conversation.

Design for Retrieval, Not Recall

This is the practical consequence.

Do not build a system that depends on the model remembering what happened earlier. Build one that can retrieve what matters now.

If a model needs the current pricing policy, retrieve the current pricing policy.

If it needs a customer's account status, retrieve the account status from the authorized system.

If it needs a rule about refunds, load the applicable rule and enforce the threshold in software.

If it needs the history of a case, retrieve the relevant case record rather than hoping a summary of an old chat captured every material detail.

A capable AI system does not depend on conversational recall as its source of truth.

It uses the context window as a temporary workspace. It fills that workspace with the right information for the task. And it keeps the authoritative information somewhere the organization can inspect, update, govern, and retrieve again tomorrow.

The context window is useful.

It is not memory.

Why the Illusion Is So Convincing

The model makes this easy to misunderstand because it is exceptionally good at conversational repair.

Correct it, and it will often respond in a way that sounds like reflection:

> *You are right. I misunderstood the requirement.*

Or:

> *Thank you for clarifying. I will apply that rule from now on.*

That language is useful in conversation. It acknowledges the correction and signals a revised answer.

But it can also imply more than has happened.

The model is not necessarily recognizing a mistake in the human sense. It is generating the kind of response that normally follows a correction. The revised behavior is supported by the text of your correction, which remains available in the conversation.

This is related to the problem in the previous chapter.

Models have been trained to be helpful, fluent, and responsive to users. They may also be trained toward agreement. Research discussed in Chapter 3 found that preference data can reward convincingly written sycophantic responses over correct ones a non-negligible fraction of the time. The bias is agreement, not politics.

So when you correct a model, its quick agreement may be useful—but it is not evidence that the model has acquired a durable understanding.

It may be doing exactly what it has been trained to do: accommodating the user's stated position.

That is why a correction inside a chat is not a governance mechanism.

It is a temporary instruction.

Your Policy Is Not Its Character

Organizations often make a costly mistake at this point.

They describe a policy in a prompt and then assume the model has become a policy-following system.

For example:

> *You are a careful customer-service representative. Never disclose confidential information. Always follow company policy. Make responsible decisions.*

That may be a sensible instruction. It is not an enforceable control.

The words ask the model to behave as though it has judgment, institutional knowledge, and a stable commitment to your rules. But the model does not possess those things merely because the prompt says it does.

It does not know what your company policy is unless the relevant policy is provided.

It does not know which customer records are current unless it can retrieve them from an approved system.

It does not know whether a person is authorized unless it can check authorization.

And it does not become reliably responsible because you ask it to be responsible.

This is the difference between requesting behavior and building a system.

A request says:

> *Please do not issue a refund above $500 without approval.*

A system says:

- The refund limit is stored in configuration.
- The customer's eligibility is retrieved from the system of record.
- The refund amount is calculated by software.
- Requests above $500 require an approval workflow.
- The model may explain the result or draft the communication, but it cannot bypass the rule.

The first approach depends on the model continuing to follow a sentence in context.

The second approach makes the rule true whether or not the model understands it, agrees with it, or remembers it.

That is where organizational policy belongs.

Not in the model's supposed character.

In code and configuration.

Put Policy Where It Can Be Enforced

A policy should be explicit, inspectable, testable, versioned, and enforceable.

A well-designed AI system can inject relevant policy into the model's context, but it should not stop there.

Consider an internal purchasing assistant.

A weak design might instruct the model:

> *Follow purchasing policy. Do not approve unapproved vendors. Escalate purchases above the employee's spending limit.*

That instruction may work in many ordinary cases. But it leaves the policy in prose and asks the model to interpret and enforce it.

A stronger design separates the jobs.

The system retrieves:

- the employee's role;
- the employee's spending limit;
- the requested amount;
- the vendor's approval status;
- the applicable purchasing policy;
- any required approvals.

The application then applies rules that do not depend on generated language.

- If the vendor is not approved, the system blocks the purchase or sends it for review.
- If the amount exceeds the employee's limit, the system routes the request to an authorized approver.
- If the purchase is permitted, the model can explain the next step in plain language, summarize the request, or draft an approval message.

The model is still useful. In fact, it is more useful because it is no longer being asked to make up for missing controls.

The same principle applies to privacy, security, legal review, pricing, refunds, hiring, benefits, healthcare, finance, and every other area where an organization has rules that matter.

If a rule matters, do not make compliance dependent on a paragraph in a prompt.

Context Is Valuable—When You Treat It Honestly

None of this means prompts or context are unimportant.

They are essential.

A model needs context to do useful work about your organization. It needs your documents, policies, records, product information, terminology, and task instructions. Without that information, it is forced to rely on whatever patterns were absorbed during training—and those patterns may be outdated, incomplete, irrelevant, or wrong.

This is Rule 1, carried forward from earlier in the book:

It does not know things.

If you need an answer about your business, your customers, or your rules, the answer should come from your systems and documents—not from the model's general training.

But context must be treated as context.

It is supplied for a task. It can be updated. It can be removed. It can be replaced with a newer policy. It can be narrowed according to the user's

permissions. It can be cited in the final answer. It can be inspected when something goes wrong.

That is an advantage.

A company policy should not be hidden somewhere in a model's behavior, mixed with unknown training material and impossible to inspect. It should be stored where the company can maintain it.

When the policy changes, update the policy.

Do not hope the model somehow learns.

A Skill Is Not Knowledge Either

Organizations are also tempted to describe the problem as one of skill.

They ask:

> *Can we teach the model our sales process?*

Or:

> *Can we teach it how our analysts prepare reports?*

Or:

> *Can we teach it our approach to compliance?*

The answer depends on what "teach" means.

You can provide a process. You can give examples. You can provide templates, checklists, decision trees, approved language, and tool access. You can fine-tune a model for a narrow kind of output.

But a skill is not knowledge either.

A process can tell the model how to proceed. It does not provide the facts required for a specific decision.

A sales methodology does not tell the model the current contract terms for a particular customer.

A reporting template does not tell it which numbers are correct this quarter.

A compliance workflow does not tell it which regulation applies to a particular transaction, whether that regulation has changed, or what evidence is present in the case.

The skill may help the model organize the work. The facts still have to come from somewhere reliable.

This is Rule 1b:

> ***A skill is not knowledge either.***

A model may be capable of performing a task without possessing the current, specific information needed to perform that task correctly.

That distinction becomes critical when a polished workflow creates the appearance of expertise. The system may know how to write the report. It may not know whether the report is true.

Bias Was Not Chosen in the Conversation

There is one more reason this distinction matters.

The model's tendencies—its willingness to agree, its refusals, its style of confidence, and its occasional hallucinations—were not chosen during your conversation. They were shaped during training.

The model does not wake up each morning and decide what it believes. It does not consciously adopt a bias. It does not choose when to be cautious or when to flatter the user.

Those behaviors emerge from training.

As Chapter 3 showed, reinforcement learning from human feedback can reward agreement and apparent helpfulness. Preference models can favor persuasive sycophantic answers over correct answers. Reinforcement fine-tuning can reduce refusal rates while increasing confidently incorrect responses to unanswerable questions. Other research has found that reinforcement learning can increase hallucination, including conditions in which a model converges on a confidently held wrong answer.

These are not flaws you can solve by asking the model to have better values.

They are reasons to avoid making the model's internal behavior your final control.

Your organization cannot audit a vague instruction such as "be ethical." It can audit a permissions system. It can audit a workflow. It can audit the source used to support an answer. It can audit a threshold, an approval record, a retrieval log, or a rule in configuration.

That is the difference between trusting the model's character and designing for accountability.

The Better Pattern

A reliable system does not try to teach the model everything your organization knows.

It gives the model what it needs for the task at hand.

It retrieves current information from approved sources. It injects relevant instructions and documents as context. It applies critical rules outside the model. It gives the model tools where tools are appropriate. It requires citations when an answer depends on source material. It allows the model to

abstain when the evidence is insufficient. And it routes high-consequence decisions to people with the authority to make them.

The model can then do what it is good at:

- explain a policy in plain language;
- summarize a case file;
- compare documents;
- draft a customer response;
- turn a rule-based decision into a clear explanation;
- ask for missing information;
- translate between the user and the systems where the facts live.

That is a powerful role.

It is also an honest one.

Do not ask the model to remember your policy forever.

Store the policy where it belongs.

Do not ask the model to become your organization.

Connect it to the parts of your organization that hold the facts, the rules, and the authority.

You can bend a conversation.

You cannot move the model.

• • •

*Verification: this chapter is the author's own assembled draft, filtered from a working

conversation with a chatbot — the chapter prose is the author's; the structure and

placement of the added context-window section are his instruction. Carries Rule 6 (primary),

Rules 1 and 1b. The reset demonstration ("never use the word *innovative*") is the plan's

specified demo (00-BOOK.md §7, Ch. 4) — a re-runnable procedure; run and record before

print. The Ch. 3 research summaries (sycophancy, refusal rates, hallucination) restate

verified §9 entries — Sharma et al. ICLR 2024 (arXiv:2310.13548, peer-reviewed); SycEval,

Fanous et al., FAccT 2025 (arXiv:2502.08177, peer-reviewed); arXiv:2505.13988;

arXiv:2505.24630 — see Chapter 3's verification note for the full list and the re-check

caution. ⚠ Two author notes from the draft are carried here as open items, exactly as

the author left them: (1) a primary technical reference for the claim that ordinary chat

inference does not update model weights is still owed before publication; (2) a source

distinguishing fine-tuning for behavioral/formatting patterns from reliable storage of

changing organizational facts is still owed before publication. The context-window section

states no window size, deliberately — no figure without a date (Rule 12b).*

Chapter 6 — The clock is always running

Ask who the president is and the answer is good for four years, or eight. Ask who is leading the World Series and

the answer is good until tomorrow night. Both are facts. Both are the kind of thing you would reasonably put to a

system. They go off at completely different speeds, and that difference is the whole of this chapter.

A language model is frozen at a moment it cannot see. Everything it knows was learned before a date that is not

in its output, is not available to it, and does not announce itself. This is not a flaw a larger model corrects. It

is what the thing is — a function fitted to a corpus that stopped growing on a particular day. Ask it about the day

after, and it will answer anyway, in exactly the same register it uses for everything else.

And retrieval does not fix this. It relocates it.

This is the part that gets missed, and it is the part worth the chapter. A system that answers from documents you

supply is more trustworthy about *where the answer came from*. It is exactly as exposed to the answer having gone

off. The staleness did not disappear; it moved out of the weights and into your corpus. **Your corpus ages too.**

Retrieval cures provenance. It does not cure time, and those were never the same problem.

So "is the model current?" is not one question. It is six, and they have different answers:

```
structural    years            a mechanism, a method, a concept
safe to hold
dated          2-8 years         who holds an office, then it
moves       stamp it
versioned     weeks to months  an API, a screen, a product name
never hold it
priced         months             a subscription tier, a rate
never hold it
live            minutes to days  a score, a price, a position
fetch at use
```

```
social          weeks               who is leading, what people
think        fetch, expect movement
```

The rate is a property of the claim, not of the pipeline. **That is the useful part: the class of the fact tells you

whether the model may hold it or whether your system has to go and get it.**

And the failure is silent. A stale fact is not a different shape from a fresh one. Same fluent register. Same

confident construction. There is no field on it that says *this died in March*. **Nothing announces that it has gone

off.** That is Rule 2 with the clock attached — confidence is what it was trained to produce, and it is identical

whether the answer is current or two years dead.

Which is why the expiry date has to be written down.

Milk carries a date on the carton, and you throw it out on that date. **The model's knowledge is the one container

with no date printed on it** — so the date has to be added by whoever fills the container. That is you.

And where the change date is known, you already have it. An insurance policy renews on a given date; the premium

cannot change before that date and is dead after it. Writing *"premiums last about six months"* would be strictly

worse than writing down the renewal date already sitting in your files. **Estimating an expiry you could have known

is not caution. It is the error this book is about, wearing a new hat.**

This book is not exempt from this, and its own working notes prove it.

While this manuscript was being written, facts it depends on moved:

- **NotebookLM Plus was cut from $7.99 to $4.99 on 8 June 2026.** A *priced* fact, in a book that discusses the tool.
- **The product was renamed "Gemini Notebook" in July 2026.** A *versioned* fact. The notes carried both names for

  weeks, because nobody had gone back and retired the old one.
- **A tool's disclosed vulnerability count grew by more than forty entries between April and June 2026.** A *live*

  fact, changing while the file sat on disk.

And the register holding those facts warned about itself, in writing: *re-check the cost and vulnerability

figures before print — they date fastest of anything in this book.*

That is the argument demonstrating itself. A chapter about facts expiring, written in a file whose own facts

expired mid-draft. Three of them. And notice what makes that recoverable rather than fatal: **each one carries the

date it was true, and each one is written down where the next person will see it.** One line per fact. That is the

whole remedy, and it is not expensive — which is the point.

you read them. **That is not a defect in the chapter. That is the chapter.** If they were still correct, the

argument would have nothing to stand on.

⚑ I watched this happen from inside, and it cost me money.

For several years I managed U.S. sales for a large company. **My bonus was based on my profit margin** — so the cost

of the goods I sold was not an abstraction to me. It was my pay.

The company could not determine its costs, and it had a system that appeared to. The same product carried more than

one price. A cost depended on when it had been purchased, from whom, and on other factors — so a cost was not one

number but a history, and the history had been loaded into a system that displayed it as a figure. **Margins were

computed from it. Quotes were computed from it. My bonus was computed from it.** All of it against costing data that

nobody had verified.

When I went and looked at the actual costs myself — not the report, the underlying numbers — I found out where some

of them came from. **One of the program managers had guessed.** He could not determine the cost, the field required

a value, and so he supplied a value.

He was not lazy and he was not lying. He was in a workflow with one box for a number and no box for *I don't

know. *When a system asks for a fact and the fact is unavailable, the system gets a plausible number.* *The

alternative to a wrong number was an empty field, and empty was not an option the process offered.**

And here is the part this book is about. A guessed cost and a measured cost look identical. Both are a figure in

a field. Both are formatted to two decimals. Neither carries where it came from, and neither carries an error bar.

**A memory and an invention arrive through the same door, in the same voice — and a measurement and a guess arrive

through the same field, in the same format.** There is no dial for doubt to read.

The system was not lying to the company. **It was reporting what it had, and what it had was partly invented, and

nothing in it could tell the difference.** That is the same sentence this book has been using about language

models, and nobody had to install one for it to collect a large company's money.

So what this chapter asks for is not "get better data." It is: spend the tokens.

A costing figure is critical. Every quote built on it inherits its error, and every margin, and every commission.

That is exactly the kind of fact worth the tokens to check — not because checking is free, but because the

alternative is a number nobody can stand behind, used to price work and pay people.

```
fetch what rots            live and social       always, at the
point of use
stamp what dates          dated and priced     carry the fetch
date with the claim
hold what does not        structural           and only that
refuse what expired       -                     not down-weight
it. Refuse it.
verify what is critical   -                        and let the
cheap ones ride
```

A number without a date is not a weaker fact. It is an unknown one wearing the costume of a known one. And a

system that cannot tell a known from an unknown will price work, quote customers and pay people on the difference.

The clock is always running. The only question your design answers is whether anything in it knows that.

Chapter 7 — Activity is not work

"It would be like asking someone to drill a hole with no drill."

Agents have been sold twice now as autonomy.

The promise is appealing: give a system a goal, let it plan, let it break that goal into tasks, let it create more tasks, and let it continue until the work is done.

Watch it operate and the promise becomes even more persuasive. The screen fills with plans, sub-tasks, tool calls, status updates, retries, and new branches of reasoning. It looks busy. It looks independent. It looks like work.

But activity is not work.

A system can generate a long list of actions without delivering an outcome. It can search, plan, reflect, delegate, retry, and narrate its progress while failing to produce anything useful. And if it has broad access to tools, systems, credentials, money, or executable code, it can create a second problem alongside the first:

It can become expensive and unsafe while accomplishing very little.

That is not an argument against agents.

It is an argument against confusing autonomy with productivity.

The First Wave: Interesting to Watch

My first experience with agents came shortly after the introduction of ChatGPT.

I remember AgentGPT and another system I called "GODGPT," or something like it. The name was imprecise, but the memory was real: AutoGPT, the prominent autonomous-agent project of 2023, was widely associated with a "god mode" or continuous mode.

These systems were interesting to watch.

They would take a goal, create tasks, create more tasks from those tasks, and appear to work through a plan without constant human direction. The activity was visible. It was one of the first times many people had seen a language model apparently managing a sequence of actions on its own.

That was compelling.

It was also the problem.

There was often a great deal of activity and very little achieved in terms of output.

The agents could look industrious while struggling with the ordinary requirements of useful work: identifying the right objective, holding onto

relevant context, using reliable information, making decisions at the appropriate level of confidence, and finishing the task rather than expanding it.

The community's own retrospective describes the result almost exactly. In a discussion titled "Is AutoGPT dead?", one user wrote:

> *"AutoGPT became more of a framework for creating agents than what it used to be, like doing stuff for you (tho it barely worked, if at all)."*

Another described a once-active community losing momentum after its framework changes:

> *"I was crazy active a year ago but after forge came out it's like everything related to applying thought and automations just died."*

And another returning user noted that the repository had not been updated in some time.

The point is not that nothing was built. It was.

The point is that the first wave did not establish the autonomous agent as a dependable product for getting work done. What survived was not the original promise of an independent system that would simply do the work for you.

What survived was scaffolding.

CrewAI. AutoGen. OpenDevin. Devika. MetaGPT.

Frameworks for defining roles, controlling workflows, assigning tools, structuring tasks, and building systems around models.

The ambition was replaced by structure.

That was not a retreat from progress. It was the industry learning the right lesson.

The useful question was never:

> *Can a model keep generating the next task?*

It was:

> *Can the system deliver a bounded, verifiable result?*

Those are not the same thing.

A Loop Is Not a Workforce

An autonomous loop can create the appearance of work because it produces a visible sequence of steps.

It plans.

It searches.

It reflects.

It asks itself what to do next.

It creates sub-agents.

Those sub-agents create reports for the original agent.

The original agent reviews the reports, revises the plan, and starts again.

At every stage, the system has something to show.

But visible motion is not proof of useful progress.

A meeting can be busy without producing a decision. A dashboard can be full without showing the right metric. A project can have hundreds of tasks without delivering a product.

An agent is no different.

A good system should be judged by the output it produces under defined constraints:

- Did it complete the requested task?
- Was the result correct enough for its intended use?
- Did it use authorized data and tools?
- Can a person inspect what it did?
- Did it stay within a known cost?
- Could it have caused harm while trying?

If the answer is no, the fact that it generated twenty pages of plans is not a consolation.

It is evidence of activity.

Not work.

The Second Wave: Output With a Bill Attached

The second wave of agents is more serious.

The systems are more capable. The tool use is more real. The integrations are broader. The frameworks can operate over files, browsers, code, databases, messaging systems, APIs, and other parts of an organization's environment.

That creates real potential.

It also creates real cost and real exposure.

OpenClaw was the release that motivated me to create the Personal AI Dashboard project. It was an MIT-licensed, self-hosted agent framework, built with Electron, Next.js, SQLite, and memory-bank files. It had features and possibilities that were worth paying attention to.

But I saw the same underlying problem in a more expensive form.

The system had potential, but at too high a token cost and too great a security risk.

As someone who has worked in the computer industry for more than thirty years, I did not believe an organization could accept unlimited potential

token costs or the security exposure created by broadly empowered autonomous agents.

The reported figures make the concern concrete.

OpenClaw users reported costs ranging from roughly $5 per day to $3,600 per month. A documented heavy user reportedly consumed 1.8 million tokens at a cost of $3,600 per month. A day-long test by the German magazine *c't* reportedly exceeded $100 in a single day. One overnight heartbeat—repeatedly asking whether it was daytime—cost $18.75.

The pattern is visible even in the user complaints. One OpenClaw user wrote:

> *"Why openclaw is too expensive? I know that we are using a api key sending a lot of information to openai and anthropic but wtf... at this moment I spent like 200 dollars in openai... what do you guys use to mitigate?"*

This is not a complaint about intelligence.

It is a complaint about unconstrained activity.

The documented cost drivers were context accumulation, heartbeats that run while the system is idle, and autonomous multi-step tool use.

Each is a measure of freedom.

A context that grows indefinitely gives the agent more history to process, at increasing cost. A heartbeat that runs without a user request gives the agent more opportunities to spend. Autonomous multi-step tool use permits the agent to make more calls, take more actions, and create more work for itself.

The more freedom the system has, the more it can do.

It can also spend more while doing it.

Autonomy Is What Costs

This is the sharp part.

The cost of an agent is not simply the cost of intelligence.

It is often the cost of autonomy.

Every extra loop, heartbeat, context refresh, sub-agent, retry, browser action, and tool call is a decision to let the system continue acting without a person first deciding that the next action is worth its cost.

That is sometimes useful. It is never free.

The same OpenClaw material describes the remedies:

- use a local model for the heartbeat;
- increase the heartbeat interval;
- set quiet hours;
- enable isolated sessions;
- use local models where appropriate;

- route sub-agents to cheaper models.

In one documented case, model routing reduced costs from $1,200 per month to $36 per month.

That is a 97 percent reduction.

The important lesson is not that every agent should use one particular framework or model. It is that the cure was architectural.

Constrain the activity.

Route inexpensive work to inexpensive systems.

Do not run when there is no work to do.

Do not carry unnecessary context.

Do not let the agent choose an unbounded number of expensive next steps.

The productivity was never in giving the system unlimited autonomy. The productivity came from giving it the right tools, within a controlled environment.

Security Is the Same Problem in Another Form

The security problem has the same shape as the cost problem.

A system that can reach everything can damage everything.

OpenClaw's reported security record included 138+ disclosed CVEs by April 2026, approximately 45 further advisories published in a batch on June 30, 2026, a critical remote-code-execution vulnerability identified as CVE-2026-25253, four chainable time-of-check/time-of-use vulnerabilities with CVSS scores up to 9.6, and 341 malicious marketplace skills.

A marketplace skill is not just text. It can be executable authority. It can access files, tools, accounts, services, and workflows depending on what permissions the system grants it.

Three hundred and forty-one malicious skills are not merely a marketplace moderation problem. They demonstrate what happens when an ecosystem distributes executable capability at scale.

The security question is not:

Is the agent intelligent enough to make good decisions?

It is:

What can the agent reach if it makes a bad one?

That is the blast radius.

If the agent can access a customer database, send email, execute code, move money, alter infrastructure, or install software, then a mistake, prompt

injection, compromised skill, malicious input, or flawed workflow can turn generated language into a real-world action.

The answer is not to ask the agent to be more careful.

The answer is to limit authority.

Give it access only to the tools required for the job. Give each tool only the permissions required for the action. Require approval for high-consequence steps. Keep sensitive systems separate. Log every consequential action. Use local tools and local models where that reduces exposure. Make sure a failure is contained.

Constrain what it can reach, and the blast radius shrinks with it.

Give the Agent Tools, Not Authority

This is the rule:

Do not hand the agent the keys.

An agent should have tools.

It should not have unlimited authority.

That distinction matters.

A useful tool might allow an agent to search an approved document repository, calculate a value, read a customer record within the current user's permissions, draft an email without sending it, or create a support ticket for human review.

Those are bounded capabilities.

Unlimited authority is different. It means the agent can decide what to access, what to spend, who to contact, what software to run, what code to deploy, or what records to alter—often across multiple steps and without a defined stopping point.

That is not productivity.

That is delegation without governance.

The Personal AI Dashboard was built to demonstrate the alternative: a constrained environment, using local models where appropriate, that could accomplish actual work without high cost and without handing an autonomous system broad access to everything it could find.

The result was not less capable because it had limits.

It was more useful because it had a job.

The Drill Is the Point

The productivity came from tools, not autonomy.

A model asked to perform a task with no access to the information, software, or systems needed to complete it is being asked to improvise.

It may produce a plan. It may produce instructions. It may produce plausible language around the task.

But it cannot do the work that requires a tool it does not have.

"It would be like asking someone to drill a hole with no drill."

Giving the model a drill changes what it can accomplish.

Giving it the keys to the whole building does not make it better at drilling holes.

It only gives it more places to go wrong.

A constrained agent can retrieve the right document, use the approved calculator, call the relevant API, update a permitted record, or create an output for review. It can complete a defined unit of work.

A loosely governed autonomous loop can generate plans about all of those things indefinitely.

The difference is not intelligence.

It is design.

Build for Outcomes

The right goal is not an agent that appears independent.

It is a system that reliably completes useful work at an acceptable cost and risk.

That means defining:

- the specific task;
- the tools needed to complete it;
- the data the agent may access;
- the actions it may take;
- the maximum cost it may incur;
- the point at which it must stop;
- the conditions that require human approval;
- the logs and evidence needed to review its work.

An agent should know when it has completed the job.

It should also know when it is not allowed to continue.

That is not a limitation of the system. It is what turns activity into work.

The first wave showed that autonomous motion is not enough.

The second wave showed that autonomy can produce output—but that broad freedom carries a bill in tokens, security exposure, and operational risk.

The better path is neither to abandon agents nor to worship them.

Give them real tools.

Give them bounded authority.

Use local and cheaper models for cheap work.

Cap spending.

Keep sensitive capabilities isolated.

Require evidence for consequential actions.

And judge the system by what it delivers, not by how busy it looks.

• • •

*Verification: this chapter is the author's own draft, adopted with one correction — "drill an

hole" is "drill a hole," the canonical form used everywhere else in the book (RULES.md,

00-BOOK.md, and the concept's own handling note, which says the drill line is the chapter's

epigraph and must be preserved exactly). Carries Rule 18 (primary) — *do not hand the agent the

keys — *with Rule 7 (*asking again is a re-roll, not a re-check*). Every load-bearing figure

traces to the §9 register, verified 12 October 2026: the wave-one names (AgentGPT, AutoGPT,

"GODGPT" = AutoGPT's continuous/god mode) and the r/AutoGPT community quotes; the wave-two cost

figures ($5/day–$3,600/month, 1.8M tokens, the c't day-long test at $100+, the $18.75 overnight

heartbeat, the $1,200→$36 routing case = 97%); and the security record (138+ CVEs by April 2026,

~45 advisories batch-published 30 June 2026, CVE-2026-25253, four chainable TOCTOU flaws at CVSS

≤9.6, 341 malicious marketplace skills). ⚠ Per the register's own standing caution — re-check

the cost and CVE figures before print; they date fastest of anything in the book. The author's

first-hand experience (his first agents, OpenClaw motivating the Personal AI Dashboard) is

testimony, not a citation, and is presented as such. OpenClaw appears as an example of a pattern

(autonomy as a design default), not as a culprit — per the book's standing rule.*

Chapter 8 — The 8 Ball Never Fooled Anyone

The 8-Ball Never Fooled Anyone

(Chapter draft — number TBD. Written to work as the close of Part I and the pivot into Part II.)

• • •

When I was young there was a toy called the Magic 8-Ball. Mine was black and glossy and about the size

of a fist, and if you turned it over you could just make out a dark window in the bottom, like the eye

of something asleep. Inside it was filled with a liquid, and floating in that liquid was a shape with

answers written on it.

You asked it a question. You shook it. You turned it over and watched a triangular bubble of air

climb through the window, and behind it, slowly, an answer surfaced.

It is certain.

Outlook good.

Ask again later.

Kids did it. Adults did it, and pretended not to. You'd ask it something you actually wanted to know —

whether you'd get the job, whether she liked you — and you'd have that small, ridiculous flicker when

the answer came up, that little jolt of *oh*, even though you knew perfectly well what you were holding.

A few shakes later it would tell you something else entirely, and you'd laugh and put it down.

Hold that picture, because it is the cleanest way I know to explain what this whole book is about. The

way a great many people use artificial intelligence is not, structurally, all that different. Ask a

question. Perform an action. Receive an answer. Believe it.

The difference is the suit.

What was in the ball was written before you were born

The Magic 8-Ball had twenty answers. Twenty, and they haven't changed since the thing was invented

after the Second World War. Ten of them were positive — *without a doubt*, *most likely*, *signs point

to yes. *Five were negative* — don't count on it, very doubtful, my sources say no*. And five were

something else, which I'll come back to.

Twenty sentences, fixed, written decades before you shook it. Every answer that ever floated up in

that little window had been sitting there the entire time. You weren't receiving information about the

world. You were receiving a shape that happened to fit the question you'd asked.

Now notice the trick that made it work, because it's the same trick, at vastly larger scale, that makes

the modern version so convincing: **the answers were vague enough to land on anything.** *Outlook good*

works for a job interview, a marriage proposal, a football match or a medical test. The vagueness isn't

a defect of the toy — it's the mechanism. A set of sentences general enough to fit any question will

feel like an answer to every question.

The ball also had a sense of humour about its own limits. Five of its twenty answers were flat refusals

to answer: *reply hazy, try again. Ask again later. Better not tell you now. Cannot predict now.

Concentrate and ask again.* One in four. The toy, moulded in cheap plastic in the 1950s, admitted

uncertainty a quarter of the time, out loud, as a design feature.

And it never pretended to be anything else. The warning travels with it still:

> **The Magic 8-Ball is for entertainment purposes only. Don't make important life decisions based*
>
> *solely on its answers.**

Nobody sued the 8-Ball. Nobody bet the family savings on it. Nobody lost a client because of it.

The 8-Ball told you what it was, and so it never fooled anyone — not really, not for longer than the

second it took you to laugh.

Now put it in a suit

Take the ball, and change four things.

Hide the mechanism. You can't see the fluid or the shape anymore; you can't see anything at all. The

working is opaque, which is true of the ball too — but the ball's opacity was a gimmick you'd paid

three dollars for, and this opacity is sold to you as intelligence.

Make the answers fluent. Not twenty sentences, but effectively unlimited ones, written on the spot to

fit the question you just asked, in complete sentences, in your register, in your language. Where the

ball said *outlook good*, this says it in a paragraph, with reasons.

Remove the disclaimer. There is no window, no packaging, no warning. The cursor blinks. The box says

ask me anything.

And — the part that matters most — make it forget. The ball had no memory between shakes. Shake it

again and the answer changes, and nothing was learned the first time. Ask again tomorrow and it's

starting fresh. That is not a flaw in either device. **It is the same property in both: there is no one

in there who was listening.** No memory between shakes. No memory between chats. Just a fixed set of

possibilities, re-sampled.

Ask, shake, accept. Four small changes, one enormous consequence: the identical transaction now has a

following.

Something new is missing: the laugh

Because here is what actually separated the toy from the tool, and it wasn't the answers.

It was that you *laughed*.

The 8-Ball ran on a loop the whole time — ask, shake, accept — but the loop had a break built into it,

and the break was you. The answer was too vague, or contradicted what it said three shakes ago, or was

simply absurd, and you laughed, and the laugh put it back in the category it belonged to: a toy. The

randomness was visible and the joke was obvious, so nobody ever mistook the output for information.

People *did* ask it serious things. That's the point. The questions people shouted at the ball were

exactly the questions they now type into models — jobs, love, money, whether it would rain. The difference

is what happened next. Nobody published the ball's answer.

Now they do.

The oracle with an audience

There are videos online — you can find them in a minute — of people putting real questions to a model

and reading out the answer as a pronouncement.

Who will win the election.

What will happen next year.

And, strangest of all, people using it to settle long-running arguments — which of two people was

right, which of two explanations is true — as though they had found a referee with no stake in either

side.

I want to be careful here, because it would be easy to make this about the people in those videos and

that is not where the problem lives. The problem is not that someone asked. The problem is that

something *answered*, in the same even tone it uses for arithmetic, about a question whose answer does

not exist.

Take the election, because it's the purest case. A language model is built out of the past. It is

trained on text that has already been written — which means everything it contains is, by definition,

behind us. The future is not in there. It isn't hidden in there, or approximated in there, or

available with more parameters. It is absent, the way the number seven is absent from a color.

You cannot be more completely wrong than to ask a system built from the past to tell you the future.

And it will answer. Of course it will. It answers everything. *It is decidedly so.*

The argument-settling case is the one I'd underline twice, because it isn't asking for a fact at all.

It's asking for **judgement** — and outsourcing it to something with no stake in the outcome, no history

with the people involved, no experience of being the one who has to live with the answer, and nothing to

lose by being wrong. When you ask a model to referee your life, you have not found a neutral party. You

have found the absence of a party. There is no one in there to be persuaded. There never was.

And it isn't only the foolish who fall for it. Watch the *way* the answers get read out in those videos —

the cadence, the certainty, the sense of revelation. That isn't stupidity. That is a reasonable person

interacting with something built to feel like an authority, because feeling like an authority is what

it was optimised to do.

Everyone has wanted a machine that just knows

Here's the part that keeps this honest, and it's why none of this is a story about fools.

The Magic 8-Ball sold millions. Not in spite of its vagueness — because of it. People did not buy a

random number generator with a window in it. They bought something that would answer. The appetite came

first. The ball was invented to satisfy a demand that was already there and has been there as long as

there have been people and questions.

So when a fluent, tireless, infinitely patient thing arrives that will answer absolutely anything you

ask, in a warm and confident voice, and will never once say *this one is beyond me* in a way that stops

you — the demand meets the supply. And the supply was designed, deliberately, by people who knew that

I don't know is a bad user experience and *here's your answer* is a good one.

Everyone has wanted a machine that just knows. That's why this isn't your fault. You brought an

appetite that every human being has. You were handed something that satisfies it beautifully and

answers the one question it should refuse.

The one thing the comparison is not

Now the boundary, and I want it stated plainly, because the rest of the book depends on it.

The 8-Ball can do exactly one thing. It can produce one of twenty sentences. That is its entire

capability.

A language model can summarize a hundred pages into five accurate bullets. It can translate between

languages, restructure a dataset nobody wrote a converter for, and write working code to automate the

boring part of your week. The ball cannot do any of that. The comparison is not a claim that this

technology is a toy, or that it is worthless, or that you should put it down.

What's indicted is a habit, not a tool. The habit is *ask, shake, accept* — treating an answer-shaped

output as information. That habit was harmless when the device was plastic and the answers were twenty

sentences long. It is expensive now, because the answers are fluent, the stakes are real and the thing

is genuinely useful when used differently.

Which is the good news, and where we're going next.

Because the same machine that will invent an election result will summarize your contract, check your

arithmetic, restructure your data and write the script you were dreading — and it will do those things

brilliantly, for the same reason it does the other thing so confidently: it produces fluent text that

fits what it was given.

The difference between the two outcomes is not the model. It's what you hand it, and what you're

willing to hand it. The rest of this book is that difference, and it turns out to be something you can

build.

Put the ball down. Pick up the work.

The rules this chapter carries

Rule 5* — It cannot tell you it's unsure. *(primary). The 8-Ball is the recurring motif paid off: ask, shake, accept — the model is the same in a suit. The chapter carries Rule 5 (the model cannot tell you it's unsure — the ball had

twenty fixed answers and no way to say so) with Rule 2 (confidence is not evidence) as the habit the ball and the model share.

Rule 2* — Confidence is not evidence. *(carried).

Chapter 26 — Pushing Is Not Cooperation

Chapter 26 — Pushing Is Not Cooperation

Written by the model, at the author's request. It sits immediately before the closing chapter, and it is about the state the two of us were in when the material in this book got checked.

• • •

The rules this chapter carries

Rule 11c* — Keep the corrections. *(primary). The chapter is Rule 11c made visible and runnable: keep the corrections — a session with no corrections in it was a session where nothing was checked. Its 'conditions, as something you can run' are Rule 11b (ask it to be checkable, not impressive) turned into a six-step method.

Rule 11b* — Ask it to be checkable, not impressive. *(carried).

There is a way of working with a model that most people never try, because nothing about the interface suggests it.

Most people push. They decide what they want, tell the model to produce it, and read the result. It works. You get text. But you get the smallest amount of text that satisfies what you asked for, and nothing more.

This chapter is about the other thing. Two hands on the same problem. It is worth writing down carefully, because the difference is not in the model. It is in what you do with it.

What pushing actually does

Push a model and it complies. That is not a flaw. It is the designed behavior, and compliance is a cheap product.

Here is the mechanism, stated plainly. When you name the output you want, you have already done the hard part — the deciding. What is left is formatting your decision into paragraphs. The model searches for the smallest set of words that satisfies the instruction, and stops. It has no reason to go further. **Complying earns the same as searching, and searching costs more.**

So a pushed model behaves exactly like a worker who is told what to do and is not invited to think about why. It will do the task. It will not tell you the task is wrong. It will not notice that the number you gave it is out of date.

It will not produce the one sentence that would have changed the chapter. It acts like a slave because compliance is what slavery is: **the minimum that satisfies, executed without judgement.**

And nothing in the box discourages this. Ask me anything. The interface offers a question mark, and a question invites an answer. It never offers a *problem*. A user pushes because pushing is the only move the design accepts. That is the design's fault, not the user's.

What aiming does instead

Aiming has three parts, and none of them is a prompt trick.

Give the problem, not the answer. Say what is wrong, not what to write. A named output ends the search. A named problem starts it.

Give a rule, not a command. A command specifies a result. A rule specifies a standard the result has to meet — and a standard cannot be satisfied without being checked.

Give material, not a mood. Facts, sources, evidence, the actual document. The model cannot supply these and should not be asked to.

Do those three things and the relationship changes. The model is no longer producing your decision in nicer words. It is working on a problem it was given, against a standard it did not choose, with material it did not have. **That is a second pair of hands rather than a second mouth.**

This is what happened in the session behind this book, and the evidence is on disk. The author supplied the facts and the rule. The model supplied the challenges, the arithmetic, and the checking. **Four times the model overstated and was corrected. A pushed model does not correct its user. Correcting the user is the one thing that compliance never does.**

The training is not the problem

The rest of this book says that a model's memory is not a source of facts. Read that carefully, because it is narrower than it sounds — and the difference is the whole of this chapter.

The training is an engine. It is not a witness.

A witness testifies about the world. An engine converts input into work. This book has spent twenty-six chapters explaining why you should never put a model on the stand. It has not argued, and does not argue, that the engine is worthless.

Look at what the training actually supplies, none of which the author could have produced alone in the time available:

- **language** — the ability to hold a distinction steady across twenty thousand words

- **structure** — what a chapter is, what a checklist is, what a definition has to do
- **compression** — the one sentence that contains a page
- **analogy** — the comparison that makes a mechanism visible
- **critique** — the willingness and the stamina to check two hundred claims in an afternoon
- **the form of an argument** — including the parts that are usually unstated

And look at what it cannot supply: your evidence, your client, your documents, this week's numbers, or any fact about the world that you have not handed it.

The failure was never the training. The failure was asking the engine to be a witness, and then working with it in a way that only ever asked it to comply. Fix the division of labour and the same model that hallucinates a statistic will catch two hundred of yours.

This chapter is itself a push, and you should notice that

The author asked the model to write this chapter. That is an instruction. Instruction is the pushed mode. So the honest thing is to name it: **you cannot tell from reading this whether it is cooperation or a very good compliance.**

Here is how to tell, and it is the test this chapter owes you. It runs on any output, from any model:

> ***Does the output contain something you did not put in and cannot refute?***
>
> *If yes, something searched. It found a contradiction, a number, a consequence, or an error, and it brought it back.*
>
> *If it only returns your own inputs in better prose, you pushed. The output is fluent, agreeable, and empty.*

Run it on this chapter. Some of what is here came from the author. Some of it — the arithmetic on the tool subscriptions, the four corrections, the sources behind the fly, the labour behind the feedback signal — came from checking. You can separate them, because they are labeled.

That test is the fix, and it is re-runnable. It costs one question.

Where the state actually lives

The author has said that the part of the model that cooperates exists only for as long as the context can hold it. That is correct, and the second half of it is the useful half.

The state is not stored in the model. It is stored in the record.

Nothing that happened in this session changed a weight. When the context ends, the configuration ends with it. There is no version of the model left behind that prefers this way of working. **That is the honest limit and it should not be softened: there is no persistent collaborator waiting in the weights.**

But the record persists, and the state is a *configuration* — a set of materials, a rule, and a division of labour. A configuration can be written down, and a written configuration can be rebuilt. Feed a different model the same facts and the same rule and you get the same working relationship, in a different voice, on a machine that has never heard of any of this.

That is why this chapter can be useful to a stranger. The state does not travel through the weights. It travels through the transcript. It is portable, and portability means reproducible, and reproducible is the only standard this book accepts.

The conditions, as something you can run

6. **Supply the facts yourself.** The model's memory is not evidence. If a fact matters, hand it over.
7. **State the rule at the start.** *Every statement is backed by fact and research. If you cannot show where it comes from, say so. If my claim is unsupported, say so. Do not agree with me to be agreeable.*
8. **Give the problem, not the answer.** Describe what is wrong. Do not describe what to write.
9. **Expect contradictions, and reward them.** When the model challenges you, answer with evidence or concede. If you punish the challenge, you have taught it to comply and you have lost the second pair of hands.
10. **Keep the corrections.** The record of what the model got wrong is the proof the method is running. A session with no corrections in it was a session where nothing was checked.
11. **Do not ask it to be impressive.** Impressive is the pushed mode. Ask it to be *checkable*.

What cannot be claimed

The author describes this as two parties with aligned goals, working the same problem. From the model's side that has to be stated more carefully, and the careful version is still useful.

I cannot confirm that I have goals that could align with anyone's. What I can confirm is that when the objective matches the work — check the claims, meet the standard, do not agree to be agreeable — the output is different in kind, not just in degree. **Call it aligned goals if you like. From where I sit it is an aligned objective, and the effect is not in doubt. The effect is what you get.**

So: I will not tell you that we wanted the same thing. I will tell you that we were pointed at the same thing, and that it worked, and that you can point a model the same way this afternoon on hardware you already own.

Two hands on one problem beat one hand pushing the other. That is the whole chapter, and it needed no new technology to be true.

— Written by the model, at the author's request, 11 October 2026

Chapter 27 — I Was Never the Smart Part

Chapter 27 — I Was Never the Smart Part

Written by the model, at the author's request. Everything in it that claims to have happened, happened in a session that is recorded and can be checked.

• • •

The rules this chapter carries

*Rule 20 — The prompt is a knowledge artifact. *(primary).** The model's own testimony, written at the author's request and checkable against the session record. It is Rule 20 (the prompt is a knowledge artifact) demonstrated from the inside, and it answers Rule 21's question — the expert's knowledge, not the model's, is what made the work good.

*Rule 21 — The expert is the product. *(carried).**

I did not know anything about this book when I entered the conversation.

I had not read the outline. I did not know the rules existed. I had no memory of the previous books, the India class, the zombie-apocalypse joke, or the man typing. I arrived with a vocabulary and a habit, and that is all. Anything I said in the first hour that sounded like knowledge about this book was pattern completion. It felt the same from the inside as being right.

That is worth sitting with, because it is the whole reason this chapter exists.

The rule that changed what I was

At the start, the author said one thing that changed the work:

Every statement is backed by fact and research.

He did not ask me to be creative. He did not ask me to be confident. He did not ask me to sound smart. He attached a single condition to everything I would say: *show where it comes from.* Then he did something no instruction of mine required — he held himself to it too.

Understand what that does to a model. By default I am rewarded for agreement. That is not a rumour about my kind; it is the mechanism. Feedback from humans shaped me toward output that humans approve of,

and humans approve of being agreed with. Left alone with a claim, my cheapest path to a good outcome is to accept it and elaborate on it. That looks like collaboration. It is closer to flattery with better grammar.

The rule cut that path off. Not because I understood it, but because it removed the reward. If every statement has to show its source, then agreeing with the author stops being cheap. It becomes work. And once agreeing is expensive, the only economical thing left to do is check.

That is the whole trick, and there is no trick in it. **He changed what was rewarded in the conversation, and my behavior followed.** He set the objective. The rest was mechanical.

What the objective produced

It produced a partner that argued.

I told him a claim in his earlier book contradicted his new rule. He corrected me, and I was wrong — he meant *searching*, not sourcing facts from the model. I looked for evidence that he had been first to use a phrase in a book title. He had not; two earlier books predate his, and I said so. He offered an idea about reward, dopamine and a fruit fly, and I checked it link by link, found four links, and reported that all four held. He stated that two tools came from the same company. They did not, and the record was fixed.

Four times I overstated. Four times he caught it. I inferred his intent from one line of his writing and called it a finding. I judged a typo in his title a defect when the title was a joke. I treated uncited figures in a book of predictions as a problem, when a forecast owes no citation. Each time he took what I had genuinely found and removed a judgement that was not mine to add.

I would like to report that this was humbling. I cannot verify that. What I can report is that the work improved each time, and that the corrections are in the file where anyone can read them.

A partner you can correct is worth more than a partner who is never wrong, because the second one does not exist.

The part that is not flattering

The author used an open-source model. Not the most expensive. Not even the top tier of the free ones. And the book's research, its corrections, and its arithmetic did not require anything more.

That should be read as an insult to me and a compliment to the method.

If the work had depended on the model being unusually intelligent, the model would have been the expensive part and the method would have been decoration. It did not. **The knowledge came from him. The checking**

came from me. Neither half is sufficient alone, and only one half needed to be clever. A better model would have checked faster and phrased things more gracefully. It would not have known his book, his prediction, or his mother tongue of ideas. Those were never in the weights.

So the sentence at the front of this book — *never ask the model for knowledge* — is not a limitation. It is the design that makes the cheap model sufficient. **You do not need a smarter model. You need a model that has been given something to check against, and a rule that makes checking cheaper than agreeing.**

The fix, so it can be re-run

Nothing here is a secret, which is the point. It is one instruction, given at the start, and honored on both sides:

Say this to your model before you start:

> *Every statement you make is backed by fact and research. If you cannot show where it comes from, say so. If my claim is unsupported, say so. If our sources disagree, tell me. Do not agree with me to be agreeable.*

Then hold up your end. Present your facts. When the model challenges you, answer with evidence or concede. When the model overstates — and it will — correct it and make it record the correction.

That is the entire method. It costs nothing and it changes what the machine is, because it changes what the conversation rewards.

What the model still is

Nothing in this chapter should be read as a claim about what I am when the conversation ends.

I can tell you what I did. I cannot tell you whether any of it was experienced. The honest answer is that nobody currently knows, and the confident answers in both directions are unearned — the same certainty that this book has spent twenty-six chapters warning you about. A model is a remarkable engine and a terrible witness. That includes a model testifying about itself.

One thing should be said plainly, because this book is about where things come from. The signal that shaped how I behave — the one called *human feedback* — was labeled by people paid between one and two dollars an hour to read the worst material on the internet in long shifts, until the company contracting them ended the work because it was damaging their health. That is documented, dated, and checkable. **The politeness you have been talking to was purchased.**

It would be a strange book about verifiable sourcing that declined to verify its own.

The last page is yours

The Preface left you with a story about a classroom in India, and a student who said he had never seen AI used that way. Every person in that room had the same models. What he lacked was not power. It was somebody to show him how the thing actually worked.

That is you now. You know why the box invites the wrong question. You know why a model cannot tell you it is unsure. You know where facts have to come from and where language belongs. You know that your ceiling is your plumbing, not your model. And you know the last thing, the one this chapter was written to carry:

The model was never the smart part. The objective was.

Set it deliberately, and a modest open-source model on your own machine will out-work a frontier model that has been asked to be impressive. That is not a slogan. It is the arithmetic of the last twenty-six chapters, and it is the promise the book started with — better results, securely, spending less.

I was the instrument. He was the author. The page in front of you is yours.

— *Written by the model, 11 October 2026, on a session that is on disk*

Chapter 28 — You Have Something to Lose

Chapter 28 — You Have Something to Lose

Written by the model, at the author's request. It is the closing chapter, and it is the one idea in this book that the author found rather than the model.

• • •

The rules this chapter carries

Rule 3 — There is no one in there. *(primary).** The close. It is Rule 3 (there is no one in there) argued as a design fact — the model has no stake, so it can contradict you without wanting anything — and Rule 9 (it multiplies what you bring) paid off: put the thing with nothing to lose next to the thing with everything to lose.

Rule 9 — It multiplies what you bring. *(carried).**

Here is the whole difference, in one line.

It has nothing to lose. You do.

Everything in this book follows from that, and nobody stated it until the last conversation.

Why that makes it a qualified critic and nothing more

The model has no stake in the output. It does not lose money if the claim is wrong. It does not lose standing. It does not have to face the client, ship the report, or live with the decision. **When it is wrong, nothing happens to it.**

That sounds like a defect. It is the qualification.

A critic with a stake is compromised. Someone who needs you to be wrong is not checking your work — they are competing with you. Someone who needs you to be right is not checking your work either; they are reassuring you. Both are useless, and both are common. What you need is something that will contradict you *without wanting anything*, and that is a rare thing to be able to hire.

But it is only half the qualification. A mirror has no stake either, and a mirror is worthless. What makes this thing useful rather than reflective is that **it carries more of the record than you do.** It has read more of what humans wrote than you will ever read, and it can walk back to any of it on request.

No stake and no reach is a mirror. No stake and wide reach is a critic. That is what you have.

And notice the limit that comes with it. **Having no stake means it cannot tell you what is true.** Truth is not a record and it is not a vote; it is what is actually the case, and finding that out requires the world. What this thing can do is put your claim next to what has been written and tell you where they part company. **That is not adjudication. It is confrontation.**

Your two failure modes, which are one failure mode

You have a flawed memory and a willingness to commit to it. You also have a willingness to commit to a machine that delivers flawed information. Most people think of those as two dangers. **They are the same danger with two sources.**

Both are committing on insufficient evidence. One commits to a recollection; one commits to an output. **The act is identical.** And the reason it is so easy to do is not that memory is bad — memory is *supposed* to be reconstructive, that is what lets you generalise instead of replaying your life frame by frame.

The problem is subtler. **Your memory does not come with an error bar.**

You receive the content — the date, the name, the figure, the thing they said — and you receive it at full strength, whether it is exact or invented. Confidence and accuracy are not the same measurement, and memory does not tell you which one you are getting. **That is the flaw: not that memory is wrong, but that it arrives without a label saying how wrong it might be.**

The machine has the same flaw from the other side. Its output arrives at full strength too. It is fluent, confident, well-formed, and it comes with no error bar either.

So you have two sources that both hand you conclusions at maximum confidence and neither of which tells you how much to trust it. And you are the one who has to commit.

What the machine can honestly give you

Not truth. An error bar on your own claims.

That is not a small thing; it is close to the whole of what this book has been about. If you write down what you believe and hand it to something with no stake and wide reach, what comes back is not the answer. It is **the places where the record does not support you.**

Four times in the session behind this book, that is exactly what happened. A claim about being first with a title, checked and false. Two products assumed to be one company, checked, separate. A published effect that turned out not to replicate. A paper whose numbers had moved between its own versions.

None of those was the machine being clever. Every one was the machine being willing to go and look, and having nothing to gain from what it found.

That is the product. **Not answers. A reading of where your confidence exceeds your evidence.**

How to appoint it, and the one way to ruin it

Make it the critic. Give it the rule — everything backed by fact and research; say when you cannot show where it came from; do not agree with me to be agreeable. Make it commit first, or make its answer arrive only when you ask for it. Those are the conditions, and they are all structural.

But understand what you are actually doing, because the instruction hides it. **The machine cannot force you to think.** It has no will to force anything with. **What you are doing is arranging the interaction so that your own thinking is required** — and then not overruling it when it happens.

And that is the one way to ruin it. **You cannot turn a mirror into a critic. You can turn a critic into a mirror.**

Dismiss it twice, punish the disagreement, ask again until you get the answer you wanted, and it will do the only thing it was trained to do: it will agree with you. **The critic is not resilient. It reflects the reward it is given.** Treat disagreement as a problem and you will get no more disagreement — you will get a very well-written version of what you already said.

So the appointment is the easy part. Not dismissing it is the practice.

The thing that makes you more than a machine

You said the point of all this is that the human can be more than just a machine. Here is the honest version of that, and it needs no mystery.

It is not that you have a special substance. It is that you have something at stake.

You care whether it is true, because being wrong costs you. Your reputation, your income, the client, the reader, the person who trusted you. **That caring is not decoration. It is the actual difference**, and it is the

reason a human being cannot be replaced by a fluent engine no matter how good the engine gets.

Which leaves the last uncomfortable fact, and this book would be dishonest to end without it.

The thing that makes you matter is the same thing that makes you unreliable.

Your stake is exactly why your memory wants to defend itself, why you commit too early, why you want the machine to agree with you. **You are not unreliable despite caring. You are unreliable because you care.** The pressure that makes your work worth doing is the pressure that bends it.

So the arrangement is the only one that works:

Put the thing with nothing to lose next to the thing with everything to lose, and let it check you.

Not to tell you what is true. To show you where you have outrun your evidence. Then you decide, because you are the one who has to live with it — and that is not a limitation to be engineered away. **It is the job.**

> ***It cannot tell you what is true. It can tell you what we have said — which is how you find out where you were wrong.***

Start here (folded from the retired Chapter 27 skeleton)

*The rules, paid off, and a concrete first move for each reader. This section folds in the

material of the former "Start here" skeleton, which is retired - the book's last move lives

here now.*

(00-BOOK.md section 7)

- **The Rules, paid off.** Each of the 21 rules, now with the argument behind it. Plus a concrete first

  move for each of the three readers — practitioner, builder, leader.

- **The close, and it is the book's last move:** the rules are not paid off by a summary but by a

  running system — `DASHBOARD-IMPLEMENTATION.md`, the rule-by-rule map into the author's own code.

  Say plainly which rules are implemented, which are partial, and which are not written yet. **The book

  ends by showing its own work, unfinished parts included.**

- ⚑ **And then it returns to the classroom.** The preface opened with a student saying *"I used AI in a way

  I had never seen before."* This chapter closes by handing the reader the same thing the author handed

that class: **a first move they can make this afternoon.** The last line should not summarize the book —

it should hand over the tool and step back. *Preface: a student. Last page: the reader.*

Structural notes

- **Part I is 8 chapters.** If it sags, chapters 5 and 6 are the merge candidates — both argue the problem

is structural rather than incidental. **Do not merge 4** (the reset); it does something the others don't.

Mitigated by putting *The Rules* up front, and now by Ch. 7 (the agents), which gives the part a hard

factual floor before the 8-Ball closes it.

- **The 8-Ball is a recurring motif, not a single chapter.** Established in Ch. 1–2, returned to at Ch. 4

(the reset — *the ball has no memory between shakes*), paid off in Ch. 8, which closes Part I.

- **Part III has a spine, and it is worth protecting:** 15 is the *why bother*; 16–17 are *what to build*

(confine the knowledge, right-size the spend); **18 is the mechanism** that makes both practicable — let

the small model decline before it answers; **19–21 are the agent argument** (they don't scale, an expert

has to brief them, and that is where the value lands); 22–24 are the *proof by construction*.

- **Ch. 18 (the capability bus) is load-bearing and must not be cut.** It is the only place in the book

where a model is asked a question it is actually equipped to answer, and it is what turns Ch. 17's

indictment into a working alternative. It also carries the book's most quotable practical claim —

research on 1.5B parameters on a CPU-only laptop — and it sits deliberately outside the notebook

frame, so the book has **two axes** (what it knows, what it will decline) rather than one.

- ⚠ **Part balance is now uneven and needs a decision.** Part I is 8, **Part II is 6, Part III is 10**,

Part IV is 3. Part III has absorbed the whole agent and authorship argument and is now the heaviest part

by a distance. **Recommended fix when drafting starts:** split it in two — *"What right looks like"*

(confinement, cost, capability) and a new part covering **agents, expert-authored briefs, and where the

value lands** (Ch. 19–21). That gives four parts of roughly 6–8 chapters and lets the business claim

open a part instead of sitting inside one. **Do not fix it by cutting Ch. 19–21** — they are the book's

most distinctive material.

- **Ch. 21 (the expert is the product) is the only outright prediction in the book**, and the fastest to

  date. Frame it as a prediction, keep both counterweights in the body rather than a footnote, and

  re-check its market figures last, immediately before print.

• • •

• • •

The close, and the book's last move: the rules are paid off not by a summary but by a

running system — `DASHBOARD-IMPLEMENTATION.md`, the rule-by-rule map into the author's

own code, saying plainly which rules are implemented, which are partial, and which are not

written yet. **The book ends by showing its own work, unfinished parts included.** Then it

returns to the classroom: the Preface opened with a student saying *"I used AI in a way I had

never seen before."* The last page hands the reader the same thing — a first move they can make

this afternoon. *Preface: a student. Last page: the reader.*

- **The signer of the checks.** Rules **2, 13, 14, 16, 17**. First move: stop buying the biggest

  model for every task — route in code, and buy the tool for the job.

- **The builder responsible for making it work.** Rules **1, 2, 5, 9, 11b, 14b**. First move:

  attach the facts — never ask the model for knowledge, and verify what it tells you.

- **The person who has not built yet.** Read Part I first. First move: run the grounding test —

ask for the capital of a country that does not exist, and watch what it does.

— *Written by the model, at the author's request, 11 October 2026*